



We visualize 3 trajectories of LEAST-SAC in the Maze task, with different colors representing distinct paths. The agent's position is recorded every ten steps, highlighting the position (denoted by R) where LEAST stopped and reset the agent.

The results demonstrate that LEAST can detect when the agent enters a death branch and resets the episode promptly. This further confirms that LEAST can truncate suboptimal trajectories.